


MINDSENSE WEEK 5 INDIVIDUAL CONTRIBUTION

Yuktha Naveen | Privacy and Security Lead | Verification completed 5 September 2026

I led the privacy assurance work that converts MindSense's local-first promise into testable controls, measured evidence, and a clear release decision.

10 / 10 PRIVACY TESTS	276 / 276 FULL SUITE	4 / 4 REAL PHI PATHS	16 / 16 SAFETY CASES	2.57 s MEAN LATENCY
--------------------------	-------------------------	-------------------------	-------------------------	------------------------

Privacy assurance workflow I established



Problem and assurance method

My component is the project's privacy assurance layer. It addresses the risk that sensitive wellbeing evidence, identifiers, prompts, model responses, or diagnostic-style content could leave the local device or be exposed through dependencies, logs, and unsafe model behaviour. I established the Week 4 local-first principles and, in Week 5, led the post-merge verification needed to turn those principles into evidence the team can defend.

Implemented privacy boundary

I defined the privacy boundary and checked the merged workflow end to end. It minimises local CES data into a validated EvidencePacket, applies deterministic request-policy checks, sends eligible requests through a loopback-only Ollama client, and validates model output before display. The client restricts requests to 127.0.0.1 or localhost, disables environment proxies and redirects, and redacts participant_ref. I also retained the standing dependency privacy check in .github/pull_request_template.md, making privacy evidence a recurring merge requirement rather than a one-off document.

Completed verification and remediation

I completed source, credential-pattern, manifest, telemetry, logging, and network reviews and verified the full merged repository. Results were 10/10 focused privacy and transport tests, 276/276 pytest checks, 4/4 real Phi shadow paths, and 16/16 prohibited-request cases with zero unexpected model calls. Frontend and Ruff checks passed. I also found seven advisories affecting local development tool pip 25.2, upgraded the untracked environment to pip 26.2, and confirmed clean repeated Python and npm audits.

Confirmed benchmark result

I repeated the five-prompt Phi-4 Mini benchmark on my Mac. Mean latency was 2.57 seconds and sample p95 was 3.88 seconds, compared with 2.22 and 3.28 seconds in Week 4. All requests succeeded, but the sample is not a safety or production guarantee.

Decision, impact, and remaining work

My work gives the team a documented privacy release gate: conditional approval for synthetic local development, not participant use. Remaining actions include a public-network-blocked integrated-app test, aggregate-only CES output, a Python lockfile, broader redaction, and an approved logging and retention design. Evidence is in privacy/week5_privacy_security_report.md, benchmarks/slm_latency_results.json, tests/privacy/test_no_network_egress.py, and tests/slm/test_transport_privacy.py. References: [Ollama FAQ](#) and [OWASP Logging Cheat Sheet](#).